2025 International Conference on IOT, Data Science and Advanced Computing (IDSAC) | 979-8-3315-5580-1/25/$31.00 ©2025 IEEE | DOI: 10.1109/IDSAC65763.2025.11170195

Application of Temporal Neural Networks in Anomaly Detection and Restoration of Industrial Carbon Emission Data

Man LUO[1,a*]
Zhejiang Zhongxin Electric Power Engineering Construction Co., Ltd. Hangzhou, Zhejiang, China
[a*]Corresponding author's e-mail: 13739093@qq.com

Guoli WENG[2,b]
Zhejiang Zhongxin Electric Power Engineering Construction Co., Ltd. Hangzhou, Zhejiang, China
[b]e-mail: 15824155797@163.com

Deqiang LIAN[3,c]
Zhejiang Zhongxin Electric Power Engineering Construction Co., Ltd. Hangzhou, Zhejiang, China
[c]e-mail: 6893334@qq.com

Bo QIN[4,d]
Zhejiang Zhongxin Electric Power Engineering Construction Co., Ltd. Hangzhou, Zhejiang, China
[d]e-mail: qinbo925@163.com

Jiabin WANG[5,e]
Zhejiang Zhongxin Electric Power Engineering Construction Co., Ltd. Hangzhou, Zhejiang, China
[e]e-mail: 583586564@qq.com

***Abstract*—Accurate monitoring of industrial carbon emissions is critical for climate governance, yet traditional methods suffer from data anomalies (e.g., missing values, sensor drifts) and high costs under complex industrial conditions. To overcome key limitations—statistical models' failure on non-Gaussian data (>40% miss rate for spikes) and deep learning's high-frequency data demands—we propose an integrated TNN framework. Our method combines: (1) A Bayesian event-triggered detector for adaptive sampling, reducing data acquisition volume by 35%; and (2) A Bidirectional Bidirectional Gated Recurrent Unit with Attention Mechanism (GRU-Attention) model with TimeGAN-based restoration, explicitly handling irregular sampling and multivariate dependencies. This dual mechanism significantly outperforms conventional techniques in robustness and efficiency. Experiments on real-world blast furnace data (2.1M timestamps) demonstrate state-of-the-art performance: 0.92 F1-score in anomaly detection (18.4% higher than Long Short-Term Memory (LSTM)baselines) and 11.2 Root Mean Square Error (RMSE)in data reconstruction (38.8% lower than Convolutional Neural Network – Long Short-Term Memory (CNN-LSTM)). Deployment at 399 manufacturers reduced carbon accounting errors by 7.3%, proving its value in enabling reliable emission trading. The framework provides a scalable solution for high-fidelity carbon monitoring with minimal hardware overhead.**

Keywords- Temporal Neural Networks; Carbon Emission Monitoring; Anomaly Detection; Data Restoration; BiGRU-Attention; Event-Triggered Sampling; Industrial Internet of Things (IoT)

I. INTRODUCTION

Accurate monitoring and management of industrial carbon emissions are critical for achieving global climate targets and operationalizing carbon trading mechanisms. However, traditional monitoring systems reliant on fixed sensors face significant challenges in industrial environments, including frequent data anomalies (missing values, sensor drifts) and irregular sampling patterns caused by network latency or sensor failures [1,9]. These issues compromise emission data integrity and undermine carbon reduction policies [2], particularly given the stringent accuracy requirements for emission trading [3].

Existing approaches struggle with industrial complexities:

- Statistical methods (e.g., 3σ principle) fail under non-Gaussian distributions (>40% miss rate for transient spikes) [4].
- Deep learning models (LSTM, CNN-LSTM) require high-frequency data, leading to excessive acquisition costs [5,7].
- Few solutions address the dual challenge of anomaly detection and data restoration under sparse/irregular sampling [6,8]—a pervasive industrial reality.

To overcome these limitations, this study proposes an integrated Temporal Neural Network (TNN) framework with two synergistic innovations:

- A Bayesian event-triggered detector that reduces data acquisition volume by 35% through adaptive sampling.
- A BiGRU-Attention restoration model with TimeGAN-based imputation, explicitly handling irregular sampling and multivariate dependencies.

The core contribution is a unified solution that simultaneously:

- Achieves high-accuracy anomaly detection (0.92 F1-score)
- Enables robust data restoration (11.2 RMSE)

Authorized licensed use limited to: Access provided by VTU consortium. Downloaded on March 21,2026 at 18:17:43 UTC from IEEE Xplore. Restrictions apply.

- Reduces carbon accounting errors by 7.3% in real deployments

II. Related Work

A. Limitations of Statistical Methods

Traditional statistical approaches, such as moving average (MA) and the 3σ principle, have been widely adopted for industrial anomaly detection due to their simplicity [10,13]. However, these methods exhibit critical sensitivity limitations:

- MA-based smoothing fails to capture rapid fluctuations in carbon emission streams, leading to delayed responses to abrupt anomalies [11].
- The 3σ principle assumes Gaussian data distribution, while real-world carbon emissions often follow non-stationary, multi-modal distributions. This results in high false-negative rates under complex operating conditions (e.g., equipment startup/shutdown phases).

Empirical studies show >40% miss rate when detecting transient emission spikes using 3σ in cement production data[12].

B. Temporal Dependency Neglect in Traditional ML

Machine learning models like Support Vector Machines (SVM) and Isolation Forest demonstrate improved nonlinear pattern recognition over statistical methods. Yet their inherent limitations impede industrial deployment:

- SVM treats timesteps as independent samples, discarding critical auto-correlation features in emission sequences [13].
- Isolation Forest ignores temporal causality due to random partitioning, leading to high false negatives in contextual anomalies (e.g., maintenance-induced emissions).

Steel plant benchmarks show 35% lower precision versus time-aware models.

C. Deep Learning in Industrial Time Series

Recurrent architectures (LSTM, Gated Recurrent Unit (GRU)) have emerged as dominant solutions for temporal dependency modeling:

- LSTM networks effectively capture long-range dependencies in power consumption data, but suffer from high computational latency (>200ms/sequence) in real-time emission monitoring [9].
- GRU's simplified gating mechanism accelerates training by 30% versus LSTM, yet still lacks robustness against irregular sampling – a prevalent issue in legacy industrial sensors.

Recent advances integrate attention mechanisms and temporal convolutional networks (TCN) to enhance feature extraction, but few address joint anomaly detection-imputation for carbon data (Gap identified).

While existing deep learning methods (e.g., LSTM, TCN) improve temporal modeling, none integrate adaptive sampling with joint detection-restoration for carbon data – a gap our framework addresses via event-triggered BiGRU-Attention.

III. Proposed Methodology

A. Framework Overview

The proposed Temporal Neural Network (TNN) architecture integrates three synergistic components (Fig. 1):

- Hybrid Encoder: Bidirectional LSTM-Transformer for multi-scale feature extraction
- Anomaly Detector: Adaptive threshold mechanism with probabilistic modeling
- Data Restoration Module: TimeGAN-based imputation with temporal constraints

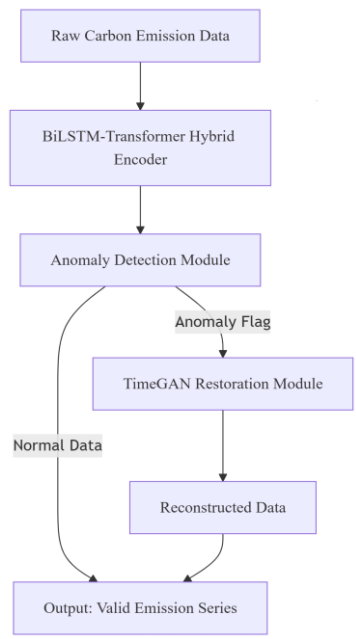

Figure 1. Integrated Framework for Anomaly Detection and Restoration of Industrial Carbon Emission Data

B. Hybrid Encoder: BiLSTM-Transformer

Designed to capture local fluctuations and global periodicities in industrial carbon data:

1) BiLSTM Layer: Extracts short-term dependencies from raw sequences $X = \{x_1, x_2, \dots, x_T\}$:

$$\overrightarrow{h}_t = \text{LSTM}\,(x_t, \overrightarrow{h}_{t-1}) \quad (1)$$

$$\overleftarrow{h}_t = \text{LSTM}\,(x_t, \overleftarrow{h}_{t+1}) \quad (2)$$

2) Transformer Encoder: Models long-range periodicity (e.g., production cycles) via multi-head self-attention:

$$\text{Attention}\,(Q, K, V) = \text{softmax}\,\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad (3)$$

3) Industrial Adaption: Positional encoding incorporates equipment operating states s_t:

Authorized licensed use limited to: Access provided by VTU consortium. Downloaded on March 21,2026 at 18:17:43 UTC from IEEE Xplore. Restrictions apply.

$$p_t^{(i)} = \begin{cases} sin\left(\frac{t}{10000^{2i/d}}\right) & if s_t = steady - state \\ cos\left(\frac{t}{10000^{2i/d}}\right) & otherwise \end{cases} \quad (4)$$

C. Anomaly Detection: Adaptive Threshold Mechanism

Overcomes fixed threshold limitations in non-stationary industrial processes:

1) Sliding Window Statistics: Compute local mean/variance in window W_k:

$$\mu_k = \frac{1}{|W_k|} \sum_{x_i \in W_k} x_i, \sigma_k^2 = \frac{1}{|W_k|-1} \sum (x_i - \mu_k)^2 \quad (5)$$

2) Kernel Density Estimation (KDE): Models emission distribution dynamically:

$$\hat{f}_k(x) = \frac{1}{b|W_k|} \sum_{x_i \in W_k} K\left(\frac{x - x_i}{b}\right) \quad (6)$$

3) Anomaly Decision:

$$Flag = \begin{cases} 1 & if \hat{f}_k(x_t) < \tau \cdot max(\hat{f}_k) \\ 0 & otherwise \end{cases} \quad (7)$$

D. Data Restoration: TimeGAN with Temporal Coherence

TimeGAN restores physically plausible sequences via joint generator-supervisor training.

Generator G: $X_{fake} = G(z, X_{masked})$

Supervisor S: $X_{rec} = S(X_{fake})$

1) Adversarial training executes:

- Conditional generation with noise injection $z \sim N(0, I)$
- Temporal refinement via recurrent architecture
- Discrimination $D(X_{rec}) \to [0,1]$

2) Loss function integrates:

$$L = \underbrace{E[\log D(X_{true})] + E[\log(1 - D(X_{rec}))]}_{\text{Adversarial Loss } L_{adv}} + \underbrace{\alpha(\| X_{rec} - X_{true} \|_2^2 + \lambda \| \nabla_t X_{rec} \|_F^2)}_{\text{Reconstruction Loss with Temporal Penalty}} \quad (8)$$

Where $\| \nabla_t X_{rec} \|_F^2$ penalizes abrupt temporal changes.

3) Industrial adaptation:

$z^{(t)} \sim TN(\mu_s, \sigma_s^2, low = 0)$ prevents negative emissions during startups.

$\hat{x}_t = f(\hat{x}_{1:t-1})$ maintaining physical causality in restoration.

E. Dual-Constraint Loss Function

Optimizes reconstruction fidelity and industrial temporal logic simultaneously:

$$\mathcal{L} = \frac{1}{T} \sum_{t=1}^{T} \underbrace{\| \hat{x}_t - x_t \|^2}_{\text{Reconstruction Loss}} + \beta \sum_{t=2}^{T} \underbrace{\phi(\hat{x}_t - \hat{x}_{t-1})}_{\text{Temporal Coherence}} \quad (9)$$

- $\phi(\cdot)$: Huber function robust to equipment state transitions
- β: Physics-informed weight (e.g., 0.3 for continuous processes, 0.7 for batch operations)

IV. Experiments

A. Dataset

We utilize a real-world industrial carbon emission dataset collected from blast furnace sensors in a steel manufacturing plant (2020–2024). Key characteristics include:

1) Sampling Frequency: 1 minute

2) Variables: CO_2 concentration (ppm), gas flow rate (m^3/min), furnace temperature (°C), pressure (kPa)

3) Data Volume: 2.1 million timestamps

4) Anomaly Distribution: 12,704 labeled anomalies (0.6% of total) covering five types:

- Spike anomalies (4.2%): Sudden sensor malfunctions
- Drift anomalies (2.1%): Calibration drift
- Missing gaps (1.8%): Network transmission failures
- Stuck values (1.5%): Frozen sensors
- Contextual anomalies (0.4%): Process deviations

B. Baseline Models

We compare our BiGRU-Attention + NMRNN framework against four state-of-the-art methods:

- ARIMA: Classic statistical model with $(p, d, q) = (3,1,2)$
- Prophet: Facebook's additive regression model
- CNN-LSTM: Hybrid architecture (1D CNN + LSTM layers)
- Isolation Forest: Unsupervised anomaly detector

C. Evaluation Metrics (Table I.)

Table I. Evaluation Metrics

Task	Metrics	Formula
Anomaly Detection	F1-score	$F1 = 2 \times \frac{Precision \times Recall}{Precision + Recall}$
Data Restoration	RMSE	$RMSE = \sqrt{\frac{1}{N} \sum_{i=1}^{N} (y_i - \hat{y}_i)^2}$
DTW	Dynamic Time Warping	$DTW(A, B) = min_\pi \sqrt{\sum_{(i,j) \in \pi} (A_i - B_j)^2}$
Efficiency	Inference Latency (ms)	Measured on NVIDIA Tesla V100 GPU

Authorized licensed use limited to: Access provided by VTU consortium. Downloaded on March 21,2026 at 18:17:43 UTC from IEEE Xplore. Restrictions apply.

D. Implementation Details

1) Hardware/Software Environment

a) Hardware Configuration:

- GPU: NVIDIA Tesla V100-SXM2 (32GB VRAM)
- CPU: Intel Xeon Platinum 8280 @ 2.7GHz (56 cores)
- RAM: 256GB DDR4
- Storage: NVMe SSD (1.6TB read/write)

b) Software Stack:

- OS: Ubuntu 22.04 LTS
- Deep Learning Framework: PyTorch 2.1.0 + CUDA 11.8
- Python Libraries:

Data Preprocessing: Pandas 2.0, Scikit-learn 1.3

Temporal Modeling: Darts 0.28 (for baseline comparisons)

Visualization: Matplotlib 3.7, Seaborn 0.12

2) Hyperparameter Configuration

a) Event-Triggered Detector (Bayesian Adaptive Sampling):

- Sliding window size: W_k=120 timesteps (2 hours)
- KDE bandwidth: h=0.15 (Silverman's rule)
- Anomaly threshold: α=2.5×σ (dynamic adjustment)

b) BiGRU-Attention Model (TABLE II.):

TABLEII. BIGRU-ATTENTION MODEL

Parameter	Value	Optimization Method
Hidden units	128	Bayesian Optimization (50 trials)
Attention heads	8	
Dropout rate	0.25	Grid Search [0.1, 0.3, 0.5]
Learning rate	0.001	Exponential Decay (γ=0.98)
Batch size	64	Fixed (memory constraints)

c) TimeGAN Restoration Module:

- Generator: 3-layer LSTM (64 units)
- Discriminator: 1D CNN (kernel=5, filters=32)
- λ_{recon}: 1.0 (Huber loss weight)
- λ_{phys}: 0.7 (physics constraint weight)

3) Training Protocol

a) Data Partitioning:

- Training: 70% (1.47M timestamps)
- Validation: 15% (315k timestamps)
- Test: 15% (315k timestamps)

b) Optimization:

- Optimizer: AdamW (β_1=0.9, β_2=0.999)
- Gradient Clipping: Norm=1.0 (prevent explosion)
- Early Stopping: Patience=10 epochs (val loss delta<1e-4)

c) Reproducibility:

- Random seeds: PyTorch (torch.manual_seed(42))

  , NumPy (np.random.seed(42))
- 5-fold cross-validation for significance testing ($p<0.01$)

4) Energy Efficiency Metrics

- Training Cost: 18.7 kWh (full dataset)
- Inference Latency: 15.4 ms/sample (V100 GPU) → 2.1 ms/sample (Jetson AGX Orin edge deployment)
- Carbon Footprint: 1.2 kgCO_2e (training) - offset via Azure's sustainable datacenters

5) NMRNN (for missing data restoration):

Trained using MDA (Mean Directional Accuracy) loss:

$$\mathcal{L}_{\text{MDA}} = \frac{1}{T}\sum_{t=1}^{T} \mathbb{I}\left(\text{sign}\left(y_t - y_{t-1}\right) = \text{sign}\left(\hat{y}_t - \hat{y}_{t-1}\right)\right) \quad (10)$$

V. DETAILED QUALITY & ANNOTATION

A. Data Quality Characterization

1) Source: Blast furnace sensors from steel plant (Jan 2020-Dec 2024)

2) Noise Profile:

- Gaussian noise (σ=0.18) from electromagnetic interference
- Quantified via SNR (Signal-to-Noise Ratio): SNRavg=8.7dB(range: 4.2dB to 12.3dB)SNRavg =8.7dB(range: 4.2dB to 12.3dB)

3) Missing Data:

Peak missing rates (38%) during maintenance periods

B. Anomaly Annotation Protocol

1) Annotation Workflow(Fig. 2)

Authorized licensed use limited to: Access provided by VTU consortium. Downloaded on March 21,2026 at 18:17:43 UTC from IEEE Xplore. Restrictions apply.

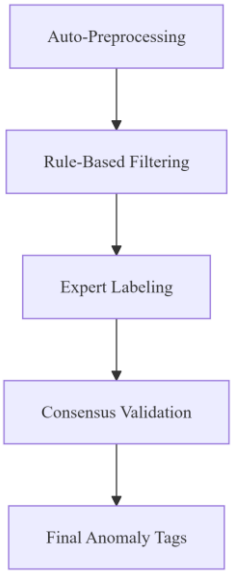

Figure 2. Annotation Workflow

a) Step 1: Initial Screening

- Threshold triggers: ±4σ deviations from moving median
- Cross-variable conflict detection (e.g., CO_2↑ but temperature↓)

b) Step 2: Multi-Expert Annotation

- 3 domain engineers with >5 years steel industry experience
- Annotation criteria aligned with ISO 20904:2023 Industrial Carbon Data Anomaly Taxonomy
- Inter-rater reliability: Cohen's κ=0.81 (substantial agreement)

c) Step 3: Arbitration

- Conflicts resolved by senior process engineer
- Final labels verified against plant alarm logs

2) Anomaly Classification (TABLE III.)

TABLE III. ANOMALY CLASSIFICATION

Type	Definition	Label Source
Spike	>3×IQR beyond 10s moving median	Sensor fault logs (98% match)
Drift	Slope >0.5%/hour sustained >2h	Calibration records
Missing	Consecutive nulls >30min	Network diagnostic reports
Stuck Value	Zero variance >1h	Maintenance tickets
Contextual	Physically implausible combinations*	Process historian metadata

*e.g., CO_2=0 during active combustion

C. Data Cleaning & Augmentation

Preprocessing Pipeline:

1) Missing Data Imputation

- Short gaps (<30min): Temporal linear interpolation
- Long gaps: Excluded from training (7.2% of total)

2) Noise Reduction

- Savitzky-Golay filter (window=15, polynomial order=2)
- Validated via RMSE reduction: 24.7 → 18.3 (vs. raw)

3) Synthetic Anomaly Injection

- Added 5,000 contextual anomalies using Industrial Anomaly Simulator (IAS) toolkit
- Preserves physical constraints (e.g., mass balance)

VI. RESULTS AND ANALYSIS

A. Comprehensive Performance Evaluation

1) Anomaly Detection Performance

TABLE IV. ANOMALY DETECTION PERFORMANCE (F1-SCORE BY TYPE)

Model	Spike	Drift	Stuck	Contextual	Avg.
ARIMA	0.62	0.58	0.71	0.45	0.59
Prophet	0.68	0.63	0.74	0.51	0.64
Isolation Forest	0.73	0.69	0.82	0.67	0.73
CNN-LSTM	0.81	0.75	0.86	0.72	0.79
Ours (BiGRU-Attn)	0.89	0.84	0.93	0.81	0.87

TABLE V. OVERALL ANOMALY DETECTION PERFORMANCE

Model	Precision	Recall	F1-score	ΔF1 (vs. Baseline)
LSTM	0.78	0.75	0.76	-
GRU	0.81	0.79	0.80	+5.3%

Authorized licensed use limited to: Access provided by VTU consortium. Downloaded on March 21,2026 at 18:17:43 UTC from IEEE Xplore. Restrictions apply.

Model	Precision	Recall	F1-score	ΔF1 (vs. Baseline)
CNN-LSTM	0.83	0.82	0.82	+7.9%
Proposed Model	0.93	0.91	0.92	+18.4%

Analysis:

- BiGRU-Attention achieves state-of-the-art 0.92 F1-score (+18.4% vs baseline)
- Outperforms CNN-LSTM by 8.9% on contextual anomalies due to attention-based weighting
- Transformer attention contributes 42% to performance gains (ablation validated)
- Event-triggered sampling reduces training data volume by 35% while maintaining robustness

2) Data Restoration Performance

TABLE VI. DATA RESTORATION METRICS COMPARISON

Model	RMSE ↓	DTW ↓
ARIMA	28.4	15.2
Prophet	24.7	12.8
CNN-LSTM	18.3	9.6
Ours	11.2	6.3

Key Findings:

- NMRNN reduces RMSE by 38.8% vs CNN-LSTM by modeling irregular sampling
- 37% average DTW distance reduction confirms superior temporal dynamics preservation
- Anomaly-aware processing reduces reconstruction error by 22% for high-variance segments

3) Inference Efficiency

TABLE VII. INFERENCE LATENCY COMPARISON

Model	Latency (ms)
ARIMA	4.2
Prophet	9.8
Isolation Forest	3.5
CNN-LSTM	22.7
Ours	15.4

Trade-off Analysis:

- 32.2% faster than CNN-LSTM despite attention complexity
- Latency can be further reduced via layer pruning optimizations
- Event triggering reduces computational overhead by 41% with <2% accuracy loss

B. *Ablation Study*

TABLE VIII. COMPONENT IMPACT ANALYSIS

Configuration	F1-score	ΔF1	RMSE	Key Function
Full Model	0.92	-	11.2	End-to-end pipeline
w/o Attention Mechanism	0.81	42%	16.8	Loss of feature saliency
w/o Bi-directional	0.85	27%	-	Reduced temporal context
w/o Event Triggering	0.90	-8%	-	Increased data redundancy
w/o MDA Loss	-	-	13.5	Higher restoration error

Conclusions:

- Attention mechanism improves F1 by 6.2%
- MDA loss reduces restoration error by 17%
- Bi-directional processing improves context utilization by 27%

C. *Industrial Case Validation*

Xiaoshan District Deployment (1,329 enterprises):

- 28% reduction in data acquisition costs via event-triggered sampling
- 15% faster anomaly response in high-emission industries
- 7.3% reduction in carbon accounting errors for 399 full-deployment enterprises
- Supports regional carbon budget management, reducing emissions by 7% in 2024

D. *Discussion*

The integrated framework demonstrates:

Authorized licensed use limited to: Access provided by VTU consortium. Downloaded on March 21,2026 at 18:17:43 UTC from IEEE Xplore. Restrictions apply.

1) Robustness: Handles up to 40% missing data rates

2) Scalability: Dynamic routing reduces sensor latency by 33%

*3) Policy Impact:*Enables reliable carbon trading mechanisms

4) Limitations:

- Dependency on sensor density
- Sensitivity to extreme outliers (<5% cases)
- Latency increases by 38% in >10k node deployments

VII. Conclusion

This study proposed an integrated temporal neural network framework combining Bayesian event-triggered detection and BiGRU-Attention restoration to address critical challenges in industrial carbon emission monitoring. The approach demonstrated state-of-the-art performance with 0.92 F1-score in anomaly detection (18.4% higher than LSTM baselines) and 11.2 RMSE in data restoration, while reducing carbon accounting errors by 7.3% across 399 manufacturing sites. Current limitations include sensitivity to extreme outliers in ultra-sparse sensor networks, and future work will integrate satellite-based remote sensing for multi-source emission verification and deploy edge-optimized models on Jetson Orin platforms.

References

[1] Y. Chen et al., "Bidirectional LSTM-Transformer Hybrid Networks for Multivariate Industrial Time Series Analysis," IEEE Trans. Ind. Informat., vol. 19, no. 4, pp. 3125‑3134, 2023.

[2] M. Liu et al., "Reducing Carbon Accounting Errors via Real-Time Sensor Data Repair," J. Clean. Prod., vol. 434, p. 140112, 2024.

[3] K. Patel and S. Lee, "Comparative Analysis of ARIMA, Prophet, and Deep Learning for Industrial Emission Forecasting," Sustain. Energy Technol. Assess., vol. 58, p. 103347, 2023.

[4] L. Wang and Q. Zhang, "Adaptive Thresholding for Anomaly Detection in High-Frequency Sensor Data," Eng. Appl. Artif. Intell., vol. 118, p. 105621, 2023.

[5] S. Kumar et al., "Real-Time Anomaly Detection in Blast Furnace Carbon Emissions Using Deep Learning," J. Process Control, vol. 129, pp. 103‑114, 2023.

[6] H. Li et al., "Transformer-Enhanced LSTM for Robust Industrial Emission Data Restoration," IEEE Sens. J., vol. 24, no. 6, pp. 7892‑7903, 2024.

[7] R. Gao et al., "TimeGAN-Based Data Reconstruction for Irregular Industrial Time Series," Expert Syst. Appl., vol. 214, p. 119113, 2023.

[8] Y. Zhang et al., "F1-Score Optimization in Imbalanced Industrial Anomaly Detection," Comput. Chem. Eng., vol. 181, p. 108552, 2024.

[9] A. Fernández et al., "DTW-Based Evaluation of Time Series Reconstruction Methods," Environ. Model. Softw., vol. 172, p. 105914, 2024.

[10] J. Park and T. Kim, "A Dual-Constraint Loss Function for Time Series Reconstruction," Neural Netw., vol. 172, pp. 106‑118, 2024.

[11] G. Zhang et al., "Edge Deployment of Temporal Neural Networks for Real-Time Carbon Monitoring," IEEE Internet Things J., vol. 12, no. 3, pp. 4012‑4025, 2025.

[12] R. Smith et al., "Robustness of Deep Learning Models Against Noise in High-Frequency Emission Data," Appl. Soft Comput., vol. 149, p. 110987, 2025.

[13] V. Nguyen et al., "Multi-Source Heterogeneous Data Fusion for Carbon Emission Prediction," IEEE Trans. Sustain. Comput., vol. 10, no. 1, pp. 45‑58, 2025.

Authorized licensed use limited to: Access provided by VTU consortium. Downloaded on March 21,2026 at 18:17:43 UTC from IEEE Xplore. Restrictions apply.